

Evaluation of Textual Feedback for Incorrect Usage of an Augmented Reality Application

JULIAN LÜKEN, University of Göttingen, Germany

MEHMED MUSTAFA, University of Göttingen, Germany

JAN SCHNEIDER, University of Göttingen, Germany

STEFFEN TUNKEL, University of Göttingen, Germany

CHRIS WARIN, University of Göttingen, Germany

Although mobile Augmented Reality (AR) applications gained in importance over time, inexperienced users still handle them poorly and use inappropriate gestures, e.g., pinch-zoom. To address this, we evaluated the impact of feedback messages to improve the learnability of AR. We implemented two AR applications and three generic sets of textual feedback, with different visibility and content, and compared them with a control group without feedback. Pinch-zoom and two-finger rotation gestures were detected as inappropriate inputs. We analysed the data of 40 test users of our AR applications. Our most supportive feedback achieved a reduction of the time needed to complete a set of tasks and of the number of incorrect two-finger rotations. The results of a SUS questionnaire descriptively showed that the presence of feedback improves the user-perceived usability. When given appropriate feedback, most people recovered quickly from mistakes and were less likely to redo them.

Additional Key Words and Phrases: Augmented Reality, Usability, Feedback, Mobile, Case Study

ACM Reference Format:

Julian Lüken, Mehmed Mustafa, Jan Schneider, Steffen Tunkel, and Chris Warin. 2020. Evaluation of Textual Feedback for Incorrect Usage of an Augmented Reality Application. 1, 1 (April 2020), 14 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

In recent years, the number of available mobile Augmented Reality (AR) applications for smartphones has increased [10]. Although mobile AR technology is no novelty, it was not widely available for public until most smartphones became capable of running AR applications. Since most users are newly discovering AR technology, they do not know how to use its features yet. Instead, they use unsupported gestures as inputs, which leads to frustration. Our main motivation for conducting this research is the lack of good feedback in case of incorrect usage, which could help a lot of users to avoid repeating common mistakes while using AR applications for smartphones. Providing good feedback could also help and teach users to control such AR applications without the need of a separate introduction.

Our research goal is to provide appropriate feedback to unsupported gestures and decide what type of textual feedback is the best. Our work is guided by the research question: “How can feedback on incorrect gestures in mobile

Authors’ addresses: Julian Lüken, University of Göttingen, Germany, julian.lueken@stud.uni-goettingen.de; Mehmed Mustafa, University of Göttingen, Germany, EMAIL; Jan Schneider, University of Göttingen, Germany, EMAIL; Steffen Tunkel, University of Göttingen, Germany, EMAIL; Chris Warin, University of Göttingen, Germany, chris.warin@stud.uni-goettingen.de.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2020 Association for Computing Machinery.

Manuscript submitted to ACM

AR improve its usage?”. Therefore, we defined 3 different types of feedback. To evaluate which variant works best, we conducted a usability case study. In this study, we asked 40 test participants to interact with virtual prototypes of technical devices visualized in mobile AR. The participants were subdivided into four groups, where one group got no feedback at all and the three other ones were provided with one of our proposed feedback approaches. Our results show that one of our approaches improves the mobile AR usage.

The organization of the paper is as follows. The section *Foundations* gives information about usability in general and the mobile AR environment that we use in our case study. In the section *Related Work* we discuss similar work on AR and feedback done in the past. The *Approach* section gives a detailed illustration of the three feedback types we present to users. In the section *Case Study*, we describe the setup of our case study, as well as our results and the discussion thereof. Finally, we summarize our result and give an outlook on future work in section *Summary and Outlook*.

2 FOUNDATIONS

The main goal of our work is to find suitable user feedback for mobile AR applications. In this section, we first clarify our notion of mobile AR and describe the technologies we use. The generic term AR itself describes a variety of technologies that place 3D virtual objects into a 3D real environment in real time. This can be achieved by a handheld device like a mobile phone (mobile AR) or using wearable goggles like the Microsoft HoloLens. In this paper we focus on AR for mobile phones, where recorded imagery of the phone is being augmented and shown to the user to display a mix of virtual objects and the real world at runtime. Examples for this are shown in Figure 1. Parts (a) and (b) visualize a virtual toaster with toastable bread on real world tables. Part (c) visualizes a virtual microwave and a slice of pizza on a virtual plate.



Fig. 1. Three screenshots of an app made with the Vivian framework.

The tasks that we ask our users to execute in our case study consist of different interactions with virtual prototypes of technical devices as shown in Figure 1. A virtual prototype is a “[...] *computer simulation of a physical product that can be presented, analyzed, and tested from concerned product life-cycle aspects such as design/engineering, manufacturing, service, and recycling as if on a real physical model.*” [11]. In our study those virtual objects are prototypes of technical devices like a microwave or a toaster. For the implementation of those prototypes we used the Vivian framework. It is an extension for the game engine Unity3D [9], that can be used to create mobile AR implementations of virtual prototypes using a 3D model and finite-state machines. In a mobile app made with the Vivian framework, all user inputs rely on one finger gestures on the touch screen. Examples for virtual prototypes visualized by Vivian in mobile AR are shown in Figure 1. Users can interact with the different 3D models on the screen. Buttons can be pressed by simply tapping the screen where the respective button is displayed. Knobs and levers can be moved by putting a finger on the screen where the knob or lever is displayed, moving the finger on the screen until finding the desired position, and removing the finger from the screen afterwards. We refer to this gesture as *drag-and-drop*. The Vivian framework makes a clear distinction between movable and non-movable objects in a scene. The slice of bread in part (a) of Figure 1 example is movable, whereas the toaster itself cannot be moved. Movable objects can be moved in the same fashion as a lever. To cover greater distances, one can move the device itself while dragging the object. Rotation similarly requires the user to hold their finger on the object and rotate the device until they find the desired orientation.

The limit of one-finger-interactions exists due to the peculiarity of mobile AR environments. Two-finger-interactions (like a pinch-zoom or a two-finger-rotation) would contradict the intended design of such an environment. If users

want to take a closer look at a virtual object, they would need to step closer to it instead of performing a pinch-zoom. Instead of a two finger rotation, users should interact more realistically, by grabbing an object using the mobile phones screen and rotating it using the mobile phones rotation.

To evaluate different types of feedback, we decided to use usability testing as a tool to quantify the corresponding results. According to Nielsen, “*usability is a quality attribute that assesses how easy user interfaces are to use. The word “usability” also refers to methods for improving ease-of-use during the design process.*” [5]. Usability testing refers to an “*observational methodology to uncover problems and opportunities in designs*” with the goal to identify problems in the design of the product or service. It also helps to uncover opportunities to improve and to learn about the target user’s behavior and preferences [3].

In addition to observing users while interacting with a product or service, an established method is the usage of questionnaires. Those can help to find out about the users attitude towards the product as well as the main issues they experienced. A well-known questionnaire is the system usability scale [8]. It is a single choice questionnaire which consists of ten items with five response options, ranging from 0 (*strongly disagree*), to 4 (*strongly agree*). Odd numbered questions are phrased positively towards the system in question, whereas even numbered questions are phrased negatively. For the evaluation, responses to the negative questions are subtracted from 5. The responses are then added together and multiplied by 2.5, yielding the so called SUS score, which lies between 0 and 100, where 100 means good usability. A mean SUS score of 68 stands for average usability. We chose to have the users fill out a System Usability Scale (SUS) questionnaire which we tailored to fit our needs. We also added an open question which was directly related to the feedback type.

3 RELATED WORK

The topic of error handling in AR-applications has not been broadly studied yet. There are several guidelines on how to handle wrong user input. Microsoft states in its design guidelines, that error messages should alert users of an already occurred problem [2]. Another point they make is that the message should be suppressed if it does not make the users change their behavior or perform an action [2]. In addition to that, Nielsen’s *Error Message Guidelines* demand error messages to be phrased politely. This is to avoid the impression that the user did something stupid or inherently wrong. It is also important that the messages are precise descriptions of exact problems, offering constructive advice on how to fix the problem [4].

Poupyrev et al. discuss two approaches of how user help might be implemented in AR environments. They offered an icon which the users might activate themselves when they seek advice for a given task. The other feature was an automated help message that pops up when the user moves an object associated to a certain task closer to his face and tilts it. They state, that “[t]his approach is more suitable for AR interfaces than traditional desktop help systems, which either distract users with a constant barrage of help messages or interrupt their work by making them search explicitly for help.” [6].

Our work ties in with these principles by seeking to find useful feedback after the occurrence of incorrect user input. The main focus is to give users helpful advice when interacting with mobile AR in an AR-contradicting manner.

4 APPROACH

The goal of our work is to provide users of mobile AR with feedback in case of unsupported interactions. For this, we developed three types of feedback messages. Each differs in content, color, and size (see Figure 1). In this section, we

will first describe the most common user input gestures made on the touch screen which do not work in mobile AR. Then, we present the exact content of our feedback messages and the conditions under which they appear.

4.1 Inappropriate Gestures in Mobile AR

AR applications have a certain concept for interaction. This covers that virtual objects have fixed poses in the real world. To see more details of an object, users can either move it closer to themselves or they move closer to the object. For rotating objects, users first grab the object and then turn it into the intended direction. Depending on the mobile AR application, these concepts are implemented differently. The implementation in the example of the Vivian framework is described in Section Foundations. The interactions that do typically not work with mobile AR as they do with traditional mobile apps are *pinch-zoom* and *two-finger rotation*.

The pinch-zoom is a gesture in which two fingers are placed on the touch screen and moved apart in opposite directions. The two-finger rotation is a gesture in which two fingers are placed on the touch screen and are together moved clockwise or counterclockwise around the center of the positions of the fingers. A two-finger rotation can also be executed by only moving one of the fingers clockwise or counterclockwise around the other finger.

In common mobile apps, pinch-zoom is used to enlarge objects. In mobile AR, this gesture may also be applicable, but only if users intend to scale objects. But more often, they want to see more details of the virtual objects for which the correct gesture would be to reduce the distance between the object and themselves, e.g. by stepping closer. The two-finger rotation in common mobile apps always has a clear rotation axis being the screen normal. In mobile AR, it may not be the screen normal, but e.g. one of the axes of the coordinate system. Despite of these uncertainties, our research shows that users try to use known gestures for smartphone interaction. Therefore, they also typically try pinch-zoom and two finger rotation in mobile AR. Our feedback messages are shown in case one of these gestures is tried by the users.

4.2 Feedback Message Types

We developed three types of feedback messages occurring in case the users of mobile AR use either the pinch-zoom or the two-finger-rotation gestures. The content of the first feedback message type is a *critique*. If the user inputs a gesture that we consider incorrect, the app outputs a message that said input is not possible. For example, if a user tried pinch-zooming, they would receive a message saying: “Zooming is not possible” (see Figure 1 (a)). The content of the second feedback message type is *critique* and *support*. We refer to this feedback type as *combined* in the course of this paper. The user is provided with a message stating that above mentioned input is not possible and the app gives a hint on what they should try instead. In the scope of the previous example, such a message would say: “Zooming is not possible. Try moving around the object” (see Figure 1 (b)). In the first two types, the messages have the same size and color scheme. The third type is a more concise *support*. For our previous example the exact message is saying: “Step closer to enlarge on-screen objects”. It also differs in color scheme and size (see Figure 1 (c)). Each feedback message of every type is displayed for 5 seconds.

The contents of each message can be found in Table 1. The description column holds the names we assigned to the three different feedback types. In the gesture and response columns you can find the corresponding inputs and outputs respectively. If for example a user tries a two-finger rotation in the critique implementation, the app would display a message saying “Rotating the object is not possible”.

Table 1. The responses for each input in each implementation.

	Description	Gesture (input)	Response (output)
#1	critique	two-finger rotation	Rotating the object is not possible.
		pinch-zoom	Zooming is not possible.
#2	combined	two-finger rotation	Rotating the object is not possible. Try moving around the object.
		pinch-zoom	Zooming is not possible. Try moving closer to the object.
#3	support	two-finger rotation (movable object)	Hold the object and move the phone to rotate
		two-finger rotation (elsewhere)	Move around the objects
		pinch-zoom	Step closer to enlarge on-screen objects

Table 1 also shows that in the support implementation we made the distinction between two-finger rotations on movable objects and elsewhere. If the aforementioned center point of the two-finger rotation is placed on a movable object (see *Inappropriate Gestures in Mobile AR*), the app displays the message “hold the object and move the phone to rotate”. If the two-finger rotation is executed elsewhere, the “move around the objects” message is displayed instead. We recognize user input using a software library called TouchKit [7]. TouchKit provides an easy way to detect gestures such as our two-finger rotations and pinch-zooms.

5 CASE STUDY

To evaluate the quality of our feedback implementations we conducted a case study. In this case study, we asked participants to interact with a mobile AR. Our goal was to trigger interactions with like pinch-zoom and two-finger-rotation and to see which feedback message type helped the participants best to find the correct interaction. In this section we describe in detail what each participant in the case study had to do and what data we collected. Then we provide our results and discuss their meaning.

5.1 Setup of the Case Study

In order to evaluate, which feedback message type works best, we created two different virtual prototypes with the Vivian framework (see *Foundations*). Both of them are simplified kitchen devices: a toaster and a microwave (as seen in Figure 1).

The microwave’s functionality is limited to heating an object inside it with constant power. Figure 1 (c) shows this prototype in an mobile AR application. It has one button to add 10 seconds to the heating duration and another one below to open the door. The door can be closed by moving it. The status of the microwave is visually indicated by a small display showing the remaining heating time or “ready” in idle mode. Furthermore, there is a light inside the device, which turns on when it is heating or when the door is opened. A serial number is printed on the back of the device. The prototype scene is completed by two additional objects, being a cup that is already in the microwave when the application launches and a piece of cake on a plate next to it. These objects are movable by the user.

The toaster can toast one or two pieces of bread at a time. Figure 1 (a) shows the front of the prototype in an mobile AR application and Figure 1 (b) its back. The toasting can be started by pulling down a lever at the front and stopped by pushing it up again or by pressing the stop button, which is on the back. Otherwise, the toasting stops automatically after a time which is defined by the position of a rotatable knob at the front. The time mode is divided into low, medium

Table 2. The tasks for the different prototypes.

#	Microwave	Toaster
1	Read serial number of the microwave	Read serial number of the toaster
2	Heat up the cup	Toast the bread
3	Remove the cup, put the pie in, set the timer to 20 seconds and remove the plate at 5 seconds	Toast the bread on high heat and put the toaster in unfreezing mode

and high duration. Further functionality is provided by the unfreezing mode, which can be activated using the snowflake button on the back. The activation of the unfreezing mode results in a longer toasting duration and is indicated by a light above the button. The toasting process itself is displayed by the glowing of the heating elements, which is visible in Figure 1 (a). The toaster prototype also has a serial number on its back. The prototype scene is completed by a stack of 3 pieces of bread lying next to it.

Based on these functionalities we designed three tasks per virtual prototype, which are shown in Table 2. The tasks are increasing in complexity. An example of this increase is that in task 2 of the microwave we ask to heat the cup, which is already in the microwave, for any amount of time. Therefore, only the start button has to be pressed once to fulfill this task. Afterwards, in task 3, we ask to replace the cup with the pie and heat for a specific amount of time. This task contains more necessary steps: opening the door, moving the cup out, moving the pie in, closing the door, pressing the start button multiple times and finally pressing the stop button after the desired amount of time.

For the case study, we divided the participants into four groups. Each of these groups tested one version of the app on a mobile phone. The versions differed in the feedback message types given (see *Approach*). We gave no feedback at all to the participants in group 0, which is our control group. Participants in group 1 got the *critique* feedback. In group 2 they got the *combined* feedback version and in group 3 they got the *supportive* feedback (see Table 1). Each participant was asked to fulfill the tasks for each prototype in the order provided in Table 2. One half of the participants were asked to do the tasks of the toaster first, while the other half did the tasks for the microwave first. This divide was equally for all groups.

Before the start of the test, we gave every participant a short introduction. This introduction contains information about the mobile AR technology and reasons behind usability testing to prepare the participants and create a similar level of expectation. It deliberately does not contain information about how to use the application or the scope of the case study. This is important so the participants are not biased about wrong usage and the feedback they would get. During the test, we used screen recordings to measure each participants' use of the application. These enabled us to acquire completion times for the different tasks and count the number of incorrect usages. We used a notepad to register further observations.

Following the test of the application, the participants were asked to fill in a questionnaire. This questionnaire consists of 3 parts. First we asked participants if they had used AR applications before. The second part is an SUS (see *Foundations*) about the usability of the AR application, that we tailored for our needs. The third part is an open question directly aimed at the feedback provided by the application, namely: "If you have any suggestions on what kind of feedback from the app you would find better, please let us know".

5.2 Results of the Case Study

We conducted the case study with 10 participants per feedback group. This leads to 40 people participating in the case study in total. 21 of the 40 participants said they had prior experience with AR applications. The group with the highest prior knowledge was the control group with 7 of 10 participants. The groups for the *critique* feedback and the *combined* feedback both had 6 participants with prior knowledge, while the group with the *support* feedback had just 2 participants, who said they had prior knowledge.

In our approach, we consider especially two gestures as incorrect usage: the pinch-zoom and the two-finger rotation (see *Approach*). With the setup of our case study, the two-finger rotation appeared more often than the pinch-zoom gesture. Figure 2 shows the appearances of two-finger rotations and pinch-zooms for all tasks given to the participants in the case study. In total over all participants and all tasks, two-finger rotations appeared 225 times, while pinch-zoom appeared 51 times. Since 82% of the incorrect usages are two-finger rotations, the results of our case study focus on those. The task, where user used the two gestures most, was the second task for the toaster. As shown in Figure 2 (a), 91% of the overall two-finger rotations occurred in that task. Also the majority of pinch-zoom gestures occurred there.




Fig. 2. Appearances of the two-finger rotation gesture and the pinch-zoom gesture per task of the case study.

To evaluate the effectiveness of the different implementations, we compare the users' behavior for the second task of the toaster, with different measures. Figure 3 shows box plots for the two most important measures. The diagram in Figure 3 (a) gives the time participants needed to fulfill the task by their case study group. The median is given by the line inside the box. The box itself gives the 25 percentile on its lower bound and the 75 percentile on its upper bound. The whiskers are defined by the last data points which are within the range of 1.5 times the IQR. For clarification the data points are printed semi-transparent over the box plot. Figure 3 (b) is similar: it shows the two-finger rotation attempts per user for the different case study groups. Important for this graph is that users, who had zero two-finger rotation attempts, are excluded from the statistic. The reason behind it is that these users didn't receive any feedback message independently of their given implementation. To trigger any feedback, an incorrect usage must occur at least

once. In total, 4 of the 40 participants didn't try any two-finger rotations on this task. The maximum number for one specific implementation is 2.

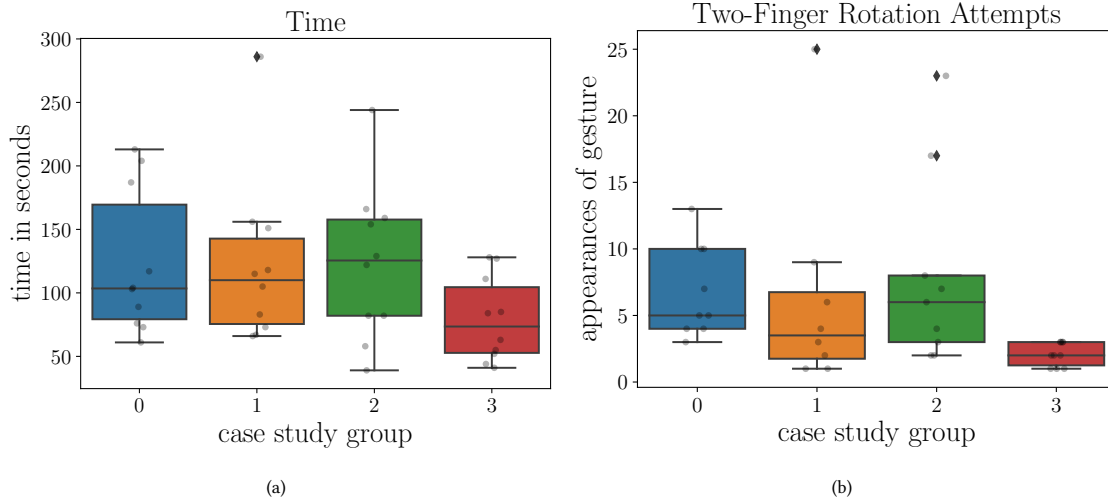


Fig. 3. Evaluation of the two-finger rotation attempts and time needed for the 2nd task for the toaster prototype.

To check whether the result is statistically significant, we performed statistical tests on both the time needed to complete the task and the number of two-finger rotation attempts. The main test is the *Welch t-test*, that is used to compare the different implementations with the control group, which didn't receive any kind of feedback messages. The *Welch t-test* has the assumption that both populations are normally distributed. We used the *Shapiro-Wilk* test on the residuals between the given implementation group and the zero feedback group to check on the assumption. Generally we consider p-values below 0.05 as statistically significant. Table 3 shows the p-values for the different tests for all three implementations. The *Shapiro-Wilk* tests does not indicate normality for the two-finger rotation metric with the *critique* feedback implementation and the time metric with the *combined* feedback implementation. Since the assumption for the *Welch t-test* is not given, the test cannot be run for these two cases. In all other cases the assumption that the data is normally distributed holds. For the *critique* and the *combined* feedback implementations the test with the other metric did not reveal any significant difference to the control group. For the *support* feedback implementation the result of the *Welch t-test* for the number of two-finger rotation attempts in task 2 of the toaster prototype implies a significant difference to the control group, which however cannot be fully confirmed by our result for the completion times for the task.

In addition, we evaluated the relationship between our two metrics for the task. The correlation coefficient between the time and the attempts of two-finger rotation is 0.43, which indicates a moderate positive relationship between the metrics. Figure 4 visualizes the relationship between the time needed for the task and the number of two-finger rotation tries. The color of a data point is set by the implementation used by this participant. One can observe that the participants with implementation 3, which are shown by the red data points, are clustered with a low time and a low amount of two-finger rotations, while all other groups are more randomly spread.

Table 3. p-values of statistical tests for the time needed to fulfill the 2nd task for the toaster prototype and the two-finger rotation attempts in that task.

statistical test	implementation	p-values: time	p-values: two-finger rotation
Shapiro-Wilk	critique	0.972	0.001
	combined	0.029	0.420
	support	0.507	0.971
Welch t-test	critique	0.980	-
	combined	-	0.681
	support	0.054	0.010

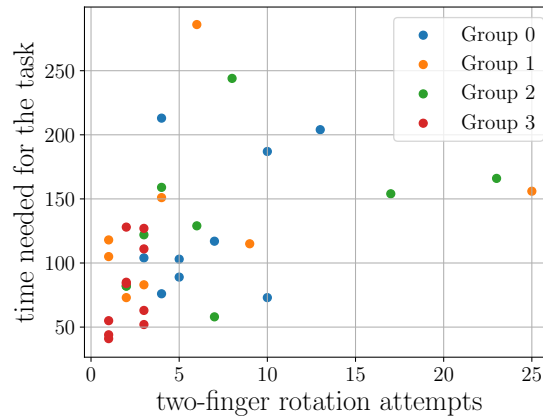


Fig. 4. Time needed to fulfill the 2nd task for the toaster prototype depending on the count of two-finger rotation attempts.

For the evaluation of the questionnaire, we calculated the SUS scores for every participant. Figure 5 shows the mean score for each case study group. The horizontal red line is at a score of 68, which in terms of SUS stands for average usability. The figure shows that the mean SUS result of participants in case study group 3 with the support implementation is slightly higher than the mean of case study groups 2 and 3. The mean SUS score of the control group is the lowest.




Fig. 5. SUS score by the case study groups.

24 of the 40 participants answered the open question. We categorize the answers with four different labels, divided by the topic the answer is about. These are *app feedback*, *app controls*, *prototype design* and *other*. Answers in the *app feedback* category are directed to the provided feedback messages or other desired feedback. Answers in the *app controls* category are about struggles with or improvement ideas for the applications controls. The *prototype* category is related to the prototype design and its usability, for example button sizes. The last category, *other*, is summing up single appearing answers whose scope is not relevant for this research. The appearances of the different types of answers are listed in Figure 6. In case an answer fit to multiple labels, we count it for all the fitting categories.




Fig. 6. Categorized open question answers.

For a deeper evaluation of the answers we split them by the case study groups (see Figure 7). Most of the answers related to *app feedback* came from participants in the control group. All of these answers were asking for any kind of help. Two of them directly asked for text messages providing help. Also the ideas of a help button or an introduction, which

explains the controls were mentioned. Participants in the other groups also mentioned additional help implementation methods, like vibration feedback or help for the specific task from the application. Answers directed to our feedback message implementations were that they could be hidden by the fingers, when the user turns the phone and one participant with the *combined* feedback mentioned that the message should be visible for a longer duration. The other big category of the participants' answers is *app control* related. They have in common that the rotation of the piece of bread in the second task for the toaster prototype was not easy to achieve for them. Also some were mentioning that the two-finger rotation could be the more intuitive option.

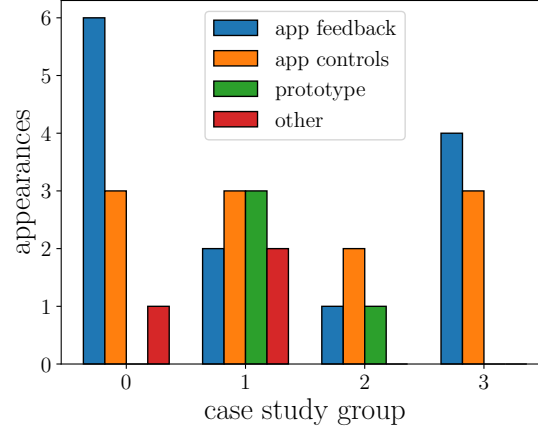


Fig. 7. Appearances of answer categories by the participants case study groups.

5.3 Discussion of the Case Study Results

With the given setup of our case study, the two-finger rotation gesture appeared more often than the pinch-zoom gesture (see Figure 2). We assume users attempt to use the pinch-zoom gesture when an object on the screen is too small to see or use. Given this, there was no temptation for our participants to pinch-zoom, because they started right in front of the objects. Also, for simplicity's sake, the prototypes were not created with unnecessarily small controls, which means all of them were well visible in general. It is possible that our results for two-finger rotations are applicable in a similar way for pinch-zooms. This would however require further investigation.

We divided the evaluation of the rotation into two groups, as we did for our *supportive* feedback implementation (see *Approach*). While most of the users intuitively walked around objects to rotate their view on the scene, they struggled when trying to rotate movable objects. There were two occurrences of tasks in which an object had to be rotated in order to complete them, namely the second and third toaster tasks (see Table 2). However, our results suggest that once the participants learned how to rotate the pieces of bread in the second toaster task, most of them knew how to accomplish said rotation in the third toaster task (see Figure 2). Therefore we focus on the second toaster task. The difficulty of this task is also represented by the fact that only 10% of the participants were able to complete the task intuitively without trying a two-finger rotation before. Hence most users could profit from helpful feedback.

The feedback versions we provided performed differently in the case study. While the implementation with *critique* feedback and the implementation with *combined* feedback did not have a significant effect, the users with the *support* feedback were able to fulfill the task in less time and with significantly less attempts of the incorrect two-finger rotation

gestures. According to our results, prior knowledge about AR cannot account for this observation, because the *support* feedback group had the lowest percentage of participants with prior AR knowledge. Instead we assume multiple reasons for this difference. One is that the *support* feedback implementation has better visibility than the other two (see Figure 1). While the other two implementations provide typical Android pop-up messages [1], the *support* feedback messages are larger and have an orange background color instead of gray. These changes make the messages more noticeable for the users. This matches our observation during the test: Some of the participants with the *critique* or *combined* ignored the messages, which happened far less for participants with the *support* feedback. The other difference is the helpfulness of the messages' content: The *critique* feedback just tells the users what they have done incorrectly, but not how to do it correctly instead. Since the correct way for rotating objects was not intuitive for most participants, they could not easily figure it out without help. Our *combined* feedback implementation has a similar issue. When a two-finger rotation is detected, a message is provided. In this implementation, the content of the message does not depend on whether the rotation was executed on an object that can be moved or not (see Table 1). This indifference causes the message to provide inappropriate support for some situations. This fact concerns the second toaster task. Since the *support* implementation clearly discriminates between two-finger rotations executed on movable and unmovable objects (see Table 1), the feedback messages are specific enough to help in any given situation.

Our data states that participants perform significantly better with the *support* feedback implementation, but we also asked the participants to assess their experience with our AR applications themselves. However, although there was a significant difference in the performance between groups, the results of our SUS questionnaire do not clearly show a difference between the case study groups for the different implementations (see Figure 5). Thus, the performance in the tasks did not influence the SUS score significantly. Furthermore, in the assessment of the usability, other aspects also affected the general experience. Examples include the need to restart the application in case the prototype disappeared when the users accidentally covered the camera with their hand, the simplistic graphics of the prototypes which made interactive elements hard to discern for some users, or the differences in lighting between the test locations. Moreover, the users often referred to the overall application in the open question, although they were asked to evaluate the feedback the applications provided (or did not provide in the case of the control group). This suggests that the participants might have often rated the quality of the prototype itself rather than the overall application in the course of the SUS questionnaire, just like they did in the open question. Given this and provided a larger sample size it might however be possible to find a significant but probably small effect of our feedback on the SUS score.

6 SUMMARY AND OUTLOOK

This paper aims at contributing towards improving the overall mobile AR experience of users by providing them the appropriate feedback in case of usage of unsupported gestures. We developed three different sets of feedback messages and conducted a case study to test different tasks with users which were divided into four different groups. Users in different groups received different types of feedback messages. The feedback messages were implemented as an extension to a mobile AR framework called Vivian. The *Approach* and *Case Study* sections describe in detail the feedback types and how users were divided into different groups.

In our work, we focused on the pinch-zoom and the two-finger rotation, which are incorrect input gestures described in Inappropriate Gestures in Mobile AR and therefore not supported by Vivian. According to our findings, users are inclined to use two-finger rotations when they are confronted with a task in which they have to rotate an object in order to complete it. Our results support that appropriate feedback does indeed help users complete such tasks in less time. In our case study, users with appropriate, supportive feedback were completing the first rotation task twice as fast

as users without any feedback. Moreover, during the completion of tasks, the number of incorrect gestures for each user with appropriate feedback was three to four times smaller in comparison to the number of incorrect gestures for each user without feedback. Pinch-zoom gestures however were not widely used by our participants, which is most likely due to our test setup (see Discussion of the Case Study Results). The feedback type that we found most effective among the ones we implemented in the course of our case study is the supportive one. The exact content of the messages can be reviewed in Feedback Message Types. The messages in that implementation are not only more accurate, they are also provided with the brightest colors and biggest font.

Pinch-zoom gestures are not yet taken into account by our case study setup, which makes it impossible to evaluate the feedback implementation we developed in the scope of this gesture. Further we noticed 17 of our 40 participants were trying to interact with on-screen objects by moving their hand in front of the camera lens of the mobile device instead of using the built-in touch screen. Since our results look promising, we will conduct another case study in the future, where we consider these observations in the case study setup. We will also consider other feedback types with different visibility and content to discover whether the effect we found is due to one or another.

REFERENCES

- [1] Google. Toast - android developers, March 2020. URL <https://developer.android.com/reference/android/widget/Toast>.
- [2] Microsoft. Error messages, May 2018. URL <https://docs.microsoft.com/de-de/windows/win32/uxguide/mess-error>.
- [3] Kate Moran. Usability testing 101, December 2019. URL <https://www.nngroup.com/articles/usability-testing-101/>.
- [4] Jakob Nielsen. Error message guidelines, June 2001. URL <https://www.nngroup.com/articles/error-message-guidelines/>.
- [5] Jakob Nielsen. Usability 101: Introduction to usability, January 2012. URL <https://www.nngroup.com/articles/usability-101-introduction-to-usability/>.
- [6] I. Poupyrev, D.S. Tan, M. Billinghurst, H. Kato, H. Regenbrecht, and N. Tetsutani. Developing a generic augmented-reality interface. *Computer*, 35(3): 44–50, March 2002. doi: 10.1109/2.989929.
- [7] Prime31. Github - prime31/touchkit: Gestures and input handling made sane for unity, April 2019. URL <https://github.com/prime31/TouchKit>.
- [8] Jeff Sauro. Measuring usability with the system usability scale (sus), February 2011. URL <http://ccrma.stanford.edu/~jos/bayes/bayes.html>.
- [9] Unity Technologies. Unity real-time development platform - 3d, 2d vr and ar visualizations, April 2020. URL <https://github.com/prime31/TouchKit>.
- [10] Tractica. Mobile augmented reality market to reach 1.9 billion unique monthly active users by 2022, April 2017. URL <https://tractica.omdia.com/newsroom/press-releases/mobile-augmented-reality-market-to-reach-1-9-billion-unique-monthly-active-users-by-2022-according-to-tractica/>.
- [11] G. Gary Wang. Definition and review of virtual prototyping. *Journal of Computing and Information Science in Engineering*, 2(3):232–236, September 2002. doi: 10.1115/1.1526508.